

AI ACADEMY, NATIONAL AI CENTER
AI-ENG-110 – SOFTWARE ENGINEERING – SPRING 2026

Final Project Report

Topic 4 - Async Research Assistant

Team: Paputu **Date:** May 23, 2026
Repo: <https://github.com/shahin1717/researchassistant>
Members: **Shahin Alakparov** (@shahin1717),
Raul Aghayev (@Raghavev17889),
Nicat Alaskarli (@NicatAlaskarli)

Executive Summary

We built an async research assistant that answers natural-language questions by querying Wikipedia, arXiv, and a configurable web-search provider in parallel, caching all source responses in SQLite, and synthesising a cited answer with a Gemini LLM. The headline concurrency result is a **4.0× speedup** (17.35s sequential → 4.37s parallel) measured across five research questions, confirming that network I/O is the dominant bottleneck and that `asyncio.gather` is the right tool. Our most important robustness story is **graceful degradation**: every source is wrapped in an individual `asyncio.timeout` and a Tenacity retry loop, so a single hanging or failing provider never blocks the answer-the pipeline synthesises from whichever sources succeeded. The most surprising lesson was that per-source timings (not just the total parallel time) are necessary for diagnosing bottlenecks: in typical runs arXiv and Tavily are 3× slower than Wikipedia, dominating the wall clock despite being parallel. If we started over we would add a **FailoverLLM** provider from day one rather than discovering the Gemini free-tier quota limit on submission day.

Contents

Executive Summary	1
1 Project Overview	3
1.1 Problem framing & scope	3
1.2 Provider choice and the AI module contract	3
2 Architecture	3
2.1 Module map	3
2.2 OOP design & key abstractions	4
2.3 Data flow across module boundaries	5
3 Concurrency & Performance	5
3.1 Why asyncio	5
3.2 Sequential-vs-concurrent benchmark	6
3.3 Rate limit & politeness	6
4 Robustness & Error Handling	6
4.1 Wrapping the AI module	6
4.2 Failure-mode analysis: Gemini 429	7
4.3 Input validation	7
5 Storage & Caching	7
5.1 Schema	7
5.2 Caching policy	8
6 Testing	8
6.1 Strategy & coverage	8
6.2 Notable tests	9
7 Deployment & Reproducibility	9
7.1 Docker image	9
7.2 Configuration	9
8 Limitations & Future Work	10
9 Tools & Acknowledgements	10
References	10
A Appendix - Reproducing the benchmark	11

1. Project Overview

1.1 Problem framing & scope

The user submits a research question via CLI or a Streamlit web interface. The system fetches relevant excerpts from three independent sources (Wikipedia, arXiv, and a web-search engine) and asks a large language model to produce a 3–6 sentence answer with inline numbered citations ([1], [2], etc.). The answer and references are printed to stdout; cost telemetry is stored in SQLite.

Inputs: a natural-language question string (1–500 characters); optional `-sources` flag restricting to a subset of providers; optional `-no-cache` flag.

Outputs: a formatted answer block with numbered references, logged per-source fetch timings, and a persistent cost entry in the telemetry database.

Entry points: `python -m src ask "..."` (CLI); `python -m src cost-report` (spend summary); `streamlit run src/app.py` (web UI).

We tightened the spec in two places: we used SQLite instead of PostgreSQL for the cache (Topic 4's TOPIC.md explicitly lists this as acceptable, and it removes infrastructure complexity from a single-container Docker image), and we did not implement the HTTP API (not required for Topic 4). We did not implement the `-since 24h` time filter on `cost-report`, which is tracked as future work.

1.2 Provider choice and the AI module contract

LLM: Google Gemini (`gemini-2.0-flash` / `gemini-3.5-flash`), selected via `LLM_PROVIDER=gemini` and `LLM_MODEL` in `.env`. The module also ships Anthropic and OpenAI adapters; swapping providers requires changing two environment variables and no source edits.

Web search: Tavily (`WEB_SEARCH_PROVIDER=tavily`). DuckDuckGo is used as a fallback when no key is configured.

We call four functions from the provided `ai/` package: `ai.fetch_wikipedia`, `ai.fetch_arxiv`, `ai.fetch_web`, and `ai.synthesize`. **We did not modify the public interface of the provided `ai/` package.**

When the provider returns a valid `AnswerWithCitations` object, the synthesizer already strips hallucinated citation indices; our SE layer additionally checks `answer.citations` is a list before iterating, and surfaces a clean `RuntimeError` (not a stack trace) when `fetch_sources` is empty.

2. Architecture

2.1 Module map

```
CLI / Streamlit UI (src/cli.py, src/app.py)
    |
    v
src/core/researcher.py          <- business logic
    |                          |
    |                          v
    |  src/concurrency/orchestrator.py  <- asyncio.gather
    |                          |
    |  [ai/ boundary]
    |  ai.fetch_wikipedia
```

```

    |   ai.fetch_arxiv
    |   ai.fetch_web
    |       |
    v       v
src/services/cache.py + ai_service.py  <- retries, TTL, rate limiter
    |
    v
src/storage/cache_store.py             <- SQLite (cache_entries, spend_log,
                                         cache_events)

```

The `ai/` package boundary is crossed only from `orchestrator.py` (for fetches) and `ai_service.py` (for synthesis). All other modules interact exclusively with SE-layer types (`Source`, `AnswerWithCitations` from `ai.schemas` are treated as read-only DTOs). Swapping providers (Anthropic → OpenAI) requires changing `LLM_PROVIDER` and `LLM_MODEL` in `.env`; zero source files change.

Module summary:

- `src/config.py` - Pydantic `Settings` reads all env vars and exposes typed fields; also initialises Python logging and forwards keys to the `ai/` package via `os.environ`.
- `src/models.py` - Frozen Pydantic models `ResearchSession` and `QueryResult` used by the SE layer; no naked dicts cross module boundaries.
- `src/concurrency/orchestrator.py` - `fetch_all()` wraps each source coroutine in `_timed()` with an `asyncio.timeout`, then fans out via `asyncio.gather`; records per-source wall-clock timings.
- `src/core/researcher.py` - Checks SQLite cache per source, calls `fetch_all` for misses, persists new results, calls `AIService.synthesize`, records cost telemetry.
- `src/services/ai_service.py` - `AIService` wraps every `ai.*` call with Tenacity exponential backoff, `asyncio.Semaphore`, per-call `asyncio.wait_for`, and an optional sliding-window `TokenBudgetLimiter`.
- `src/services/cache.py` - `QueryCache` normalises keys (`.strip().lower()`), delegates to `CacheStore`, and logs hit/miss events.
- `src/storage/cache_store.py` - Thread-safe SQLite with three tables: `cache_entries`, `spend_log`, `cache_events`. Exposes telemetry helpers (`spend_breakdown`, `most_expensive_queries`).
- `src/cli.py` - argparse CLI with `ask` and `cost-report` subcommands; validates question length and source aliases before any network call.
- `src/app.py` - Streamlit UI with a Research tab and a Developer Dashboard (cache stats, spend breakdown, expensive queries).

2.2 OOP design & key abstractions

The most important OOP decision in our SE layer is `TokenBudgetLimiter`, which uses **composition** inside `AIService`: the service holds a `TokenBudgetLimiter | None` instance and delegates the sliding-window bookkeeping to it.

```

1 class TokenBudgetLimiter:
2     """Sliding-window token budget (TPM) limiter."""
3
4     def __init__(self, tokens_per_window: int, *,
5                 window_seconds: float = 60.0,
6                 sleeper: Callable[[float], Awaitable[None]] = asyncio.sleep,
7                 clock: Callable[[], float] = time.monotonic) -> None:

```

```

8         self._tokens_per_window = tokens_per_window
9         self._window_seconds = window_seconds
10        self._sleep = sleeper
11        self._clock = clock
12        self._lock = asyncio.Lock()
13        self._events: deque[_TokenUsage] = deque()
14
15        async def acquire(self, tokens: int) -> float:
16            """Block until tokens are available. Returns seconds waited."""
17            ...
18
19    class AIService:
20        def __init__(self, *, token_budget_tpm: int | None = None, ...) -> None:
21            self._rate_limiter = (
22                None if token_budget_tpm <= 0
23                else TokenBudgetLimiter(token_budget_tpm, ...)
24            )

```

We chose composition over inheritance because `TokenBudgetLimiter` has no natural “is-a” relationship with `AIService` - it is a self-contained policy object. Making it injectable (via constructor parameters) also enables the test suite to pass a fake `sleeper` and `clock` without patching globals, which made the token-budget tests deterministic.

The provided `ai/` package already demonstrates the ABC/factory pattern (`LLMProvider` → `AnthropicProvider`, `OpenAIProvider`, `GeminiProvider`). We reference rather than replicate that pattern; our own example of a polymorphic design is `QueryCache`, which accepts any `CacheStore` instance, making it straightforward to substitute an in-memory store in tests.

2.3 Data flow across module boundaries

Our SE-layer dataclasses:

- `ResearchSession` - frozen Pydantic model carrying the validated question, source list, and cache config for a single request.
- `QueryResult` - frozen Pydantic model returned by `QueryCache.get/set`; carries the canonical key, raw JSON payload, and TTL metadata.

No naked dictionaries cross module boundaries. The `Source` and `AnswerWithCitations` types from `ai.schemas` are used as read-only DTOs within the SE layer without wrapping, because they are already well-typed frozen Pydantic models. We wrapped the cache payload (`QueryResult`) because the SE layer needed to add `cached`, `expires_at`, and `canonical_query` fields that the `ai/` schema does not provide.

3. Concurrency & Performance

3.1 Why asyncio

We used `asyncio` + `httpx` for all three source fetchers. The bottleneck is network I/O: each fetch is a single HTTP round-trip to an external API (Wikipedia REST, arXiv Atom, Tavily). Threads would work but carry heavier overhead and do not compose naturally with `asyncio.timeout`. `ProcessPoolExecutor` would be wrong because there is no CPU-bound computation in the fetch path.

Parallelism is bounded by `asyncio.Semaphore(MAX_PARALLEL)` (default 10) inside `AIService._run_async_call`. We chose 10 because each query generates at most three concurrent source calls, so the semaphore is effectively never the limiting factor for a single question; it protects against accidental misuse if the CLI is driven with a loop. With `Semaphore(1)` the system degrades to sequential; with

`Semaphore(100)` on a high-QPS workload it would exhaust provider rate limits and trigger retries.

A single shared `httpx.AsyncClient` is passed to all three fetchers per request, enabling HTTP connection reuse across sources.

3.2 Sequential-vs-concurrent benchmark

Workload	N	Sequential (s)	Concurrent (s)	Speedup
5 research questions (wiki + arxiv + web)	5	17.35	4.37	4.0×

Machine: MacBook Pro M1, Python 3.11, caches cleared, `scripts/bench.py`.

The theoretical maximum speedup for three equal-latency sources is 3×; we exceed this because Wikipedia is significantly faster than arXiv and Tavily. In the parallel case the wall clock equals the slowest source (~1.5s for Tavily), while the sequential case sums all three (wiki 0.5 + arxiv 1.4 + web 1.5 = 3.4s, multiplied across 5 questions). The new bottleneck is the LLM synthesis call (~5s per question), which is serial and dominates total end-to-end latency. The per-source timing breakdown is now visible in the logs:

```
INFO source_fetched source=wiki elapsed_s=0.48 count=3
INFO source_fetched source=arxiv elapsed_s=1.37 count=3
INFO source_fetched source=web elapsed_s=1.54 count=3
INFO fetch_all_complete timings={...total_parallel: 1.54} total_sources=9
```

3.3 Rate limit & politeness

Every `ai.*` call is wrapped in a Tenacity retry with exponential backoff (`wait_exponential(multiplier=0.5, min=0.5, max=5.0)`, 3 attempts). On HTTP 429, the `ProviderError` is caught and retried with back-off.

For token-per-minute enforcement, `TokenBudgetLimiter` maintains a sliding deque of (`timestamp`, `tokens`) events. Before each AI call it estimates the token count from argument character lengths (1 token ≈ 4 characters), prunes events older than 60s, and sleeps until headroom opens if the budget would be exceeded. The TPM limit is configurable via `TOKEN_BUDGET_TPM`; it defaults to 0 (disabled) so the free-tier demo does not stall. We observed a 429 from Gemini during development when the free-tier daily quota was exhausted; the logs showed three `ai_call_retry` warnings followed by `ai_call_failed`, and the CLI printed a clean `Error: Gemini call failed: 429 RESOURCE_EXHAUSTED` message.

4. Robustness & Error Handling

4.1 Wrapping the AI module

Every call to `ai.*` goes through `AIService._run_async_call`, which:

1. Acquires the token budget (sleeps if TPM would be exceeded).
2. Logs `ai_call_start` with `operation` and `timeout_seconds`.
3. Acquires `asyncio.Semaphore` to bound concurrency.
4. Calls `asyncio.wait_for(func(...), timeout=timeout_seconds)`.
5. On success, logs `ai_call_success` with `duration_ms`.

6. On `ProviderError` or `TimeoutError`, Tenacity retries up to 3 times with exponential back-off, logging `ai_call_retry` each time.
7. After all retries exhausted, logs `ai_call_failed` and re-raises.

Each source additionally sits inside a per-task `asyncio.timeout(per_source_timeout_seconds)` in `orchestrator.py`, so a slow Wikipedia response cannot hold arXiv and Tavily hostage regardless of how many retries are attempted. The synthesis call uses a separate 30s timeout.

Input validation before calling the AI module: question is stripped and checked non-empty and ≤ 500 characters; `sources` must contain at least one recognised alias (`wiki`, `arxiv`, `web`); empty `fetches_sources` raises `RuntimeError` before the synthesis call is even attempted.

4.2 Failure-mode analysis: Gemini 429

How we injected it: ran repeated queries until the Gemini free-tier daily quota was exhausted.

What the system did: Tenacity fired three attempts at 0.5s, 1s, and 2s back-off. All three returned 429. After the third attempt the exception propagated up through `researcher.py` to the CLI error handler.

What the user saw:

```
Error: Gemini call failed: 429 RESOURCE_EXHAUSTED. ...
* Quota exceeded for metric: .../generate_content_free_tier_requests
Please retry in 22s.
```

What the logs showed:

```
INFO ai_call_start operation=synthesize timeout_seconds=30.0
WARN ai_call_retry attempt=1 error=ProviderError
WARN ai_call_retry attempt=2 error=ProviderError
ERROR ai_call_failed operation=synthesize duration_ms=6482.3
ERROR research_failed error=Gemini call failed: 429 ...
```

The log was sufficient to diagnose the root cause without re-running the code. The fix is to set `LLM_PROVIDER=anthropic` in `.env` and supply an Anthropic key.

4.3 Input validation

- **CLI:** `validate_question()` rejects empty strings (*“Question cannot be empty.”*) and strings > 500 characters (*“Question must be 500 characters or fewer.”*). `parse_sources()` rejects unknown source aliases with a list of allowed values.
- **Env:** Pydantic `Settings` with typed fields; missing required keys produce a `ValidationError` on startup with a clear field name.
- **Cache layer:** `QueryCache._normalize_source()` raises `ValueError` on empty source strings. `CacheStore.record_cache_event()` raises `ValueError` if event is not `'hit'` or `'miss'`.
- **Models:** `ResearchSession` and `QueryResult` use `ConfigDict(extra="forbid")` and field validators to reject blank strings.

5. Storage & Caching

5.1 Schema

```

CREATE TABLE IF NOT EXISTS cache_entries (
    source          TEXT NOT NULL,
    canonical_query TEXT NOT NULL,
    response_json   TEXT NOT NULL,  -- JSON array of ai.schemas.Source
                                dicts
    created_at      REAL NOT NULL,  -- Unix timestamp (time.time())
    PRIMARY KEY (source, canonical_query)
);

CREATE TABLE IF NOT EXISTS spend_log (
    id              INTEGER PRIMARY KEY AUTOINCREMENT,
    source          TEXT NOT NULL,  -- provider name (e.g. "synthesis")
    canonical_query TEXT NOT NULL,
    cost_usd        REAL NOT NULL,
    created_at      REAL NOT NULL
);

CREATE TABLE IF NOT EXISTS cache_events (
    id              INTEGER PRIMARY KEY AUTOINCREMENT,
    source          TEXT NOT NULL,
    canonical_query TEXT NOT NULL,
    event           TEXT NOT NULL CHECK(event IN ('hit', 'miss')),
    created_at      REAL NOT NULL
);

```

`response_json` stores a JSON-serialised `list[Source.model_dump()]`. Switching embedding providers or source schemas requires a cache flush, not a migration. We chose text/JSON over binary blobs because the payloads are small and human-readable, which eases debugging. `CacheStore` is thread-safe via `threading.RLock`.

5.2 Caching policy

We cache per `(source, canonical_query)` pair with a configurable TTL (default 24h). The canonical query is `query.strip().lower()`, so "WHAT IS PHOTOSYNTHESIS?" and "what is photosynthesis" hit the same row. `cleanup_expired()` is called on every `cache.get()` to remove stale rows before the lookup.

In a typical demo run (5 questions, first time), all 15 source lookups are misses. On a second run without `-no-cache` all 15 are hits, reducing total latency from ≈ 35 s to ≈ 25 s (only synthesis calls remain). The hit rate from `artefacts/demo_output.txt` (second run): $15/15 = 100\%$.

The main staleness risk is that web-search results change over time. A news-related question cached yesterday may omit today's events. Production hardening would require shorter TTLs for web sources and a freshness indicator in the answer.

6. Testing

6.1 Strategy & coverage

Suite	Coverage
src/ (SE layer, "pytest -cov src/")	86%
Provided ai/ smoke tests	passing

All 51 tests run offline. The `ai.*` functions are mocked via `unittest.mock.AsyncMock` and custom `FakeAIService` classes in `conftest.py`. HTTP is not directly mocked because all HTTP passes through the provided `ai/` package, which is itself mocked at the `ai.fetch_*` boundary. The three

main uncovered areas are `src/app.py` (Streamlit UI, not unit-testable), `src/cli_demo.py` (demo script), and some late error-path branches in `ai_service.py`.

Provided smoke tests (`tests/test_ai_smoke.py`) pass unchanged.

6.2 Notable tests

(i) **Happy-path end-to-end** (`test_se_layer.py::test_researcher_returns_cited_answer`): mocks all three sources returning one `Source` each and verifies the final `AnswerWithCitations` has a non-empty answer and at least one citation. Exercises the full stack from `researcher.research()` down to the mock synthesizer.

(ii) **Concurrency / graceful degradation** (`test_se_layer.py::test_orchestrator_per_task_timeout_isolation`): patches `arXiv` to `asyncio.sleep(9999)` and sets `per_source_timeout_seconds = 0.05`. Asserts that Wikipedia results still appear in the output and that the total wall-clock is under 1s. This test proved the per-task `asyncio.timeout` was necessary-before adding it, the test hung for the full 9999 seconds.

(iii) **Error-path coverage** (`test_ai_service.py::test_token_budget_limiter_waits_when_budget_exhausted`): fills the token budget to capacity, then calls `acquire(1)` and asserts the limiter slept at least once. Caught a bug where `_seconds_until_available` returned a negative value for an already-expired window, causing an infinite loop.

7. Deployment & Reproducibility

7.1 Docker image

```
# Build
docker build -t research-assistant .

# Run CLI
docker run --env-file .env research-assistant \
  python -m src ask "What is photosynthesis?"

# Run Streamlit UI
docker run --env-file .env -p 8501:8501 research-assistant \
  streamlit run src/app.py --server.address 0.0.0.0
```

The Dockerfile uses a **two-stage build**: a `python:3.11-slim` builder installs dependencies into `/opt/venv`; the runtime stage copies only the venv and application code. The image runs as a non-root appuser. Python version: 3.11-slim.

7.2 Configuration

Variable	Required	Missing behaviour
LLM_PROVIDER	no (default: <code>gemini</code>)	Uses Gemini
LLM_MODEL	no (default: <code>gemini-2.0-flash</code>)	Uses default model
GOOGLE_API_KEY	if Gemini	API calls fail with 401
ANTHROPIC_API_KEY	if Anthropic	API calls fail with 401
OPENAI_API_KEY	if OpenAI	API calls fail with 401
WEB_SEARCH_PROVIDER	no (default: <code>tavily</code>)	Falls back to DuckDuckGo
TAVILY_API_KEY	if Tavily	Web search disabled or DDG used
CACHE_DIR	no (default: <code>./.cache</code>)	Creates directory automatically
CACHE_TTL_SECONDS	no (default: 86400)	24-hour TTL
PER_SOURCE_TIMEOUT_SECONDS	no (default: 10)	10-second per-source cap
MAX_PARALLEL	no (default: 10)	Semaphore bound of 10
LOG_LEVEL	no (default: <code>INFO</code>)	INFO logging
TOKEN_BUDGET_TPM	no (default: 0)	Rate limiter disabled

8. Limitations & Future Work

- **Free-tier LLM quota.** The system uses Gemini’s free tier by default, which has a per-day request cap. On submission day we hit this limit. Production deployment requires a paid API key or a configured fallback provider. A `FailoverLLM` wrapper (bonus +3) would have resolved this automatically.
- **Wikipedia often returns zero results.** The Wikipedia REST API regularly returns an empty list for technical queries. We cache the empty result, which is correct but means subsequent runs also show 0 Wikipedia sources. A smarter retry with a shorter query would help.
- **Commit imbalance.** Shahin has ~82% of commits (combining both git identities) due to being the integration lead; Raul and Nicat contributed core modules (services layer and CLI respectively) but their commits are fewer. This is documented in the contribution statement.
- **No -since filter on cost-report.** The spend log stores timestamps but the CLI does not yet expose time-range filtering. All spend since the database was created is reported.
- **SQLite in a multi-process deployment.** The current `RLock` strategy is correct for a single-process async application. A multi-worker deployment (e.g., Gunicorn + Uvicorn workers) would require migrating to PostgreSQL or a shared Redis cache.

9. Tools & Acknowledgements

AI-assisted contributions:

- `src/concurrency/orchestrator.py` - Claude (Anthropic) suggested the `_timed()` coroutine wrapper pattern and the placement of `asyncio.timeout()` per task. The team reviewed and understands every line; specifically, that `_timed` catches all exceptions so `asyncio.gather` needs no `return_exceptions=True`.
- `tests/test_storage.py` and `tests/test_researcher.py` - Claude scaffolded the initial test structure; the team added domain-specific assertions and corrected the `conftest.py` fixture references.
- `scripts/demo.py` - Claude drafted the script; the team integrated it with the actual `format_answer` helper and verified it against the live demo.
- `src/config.py` (`_ExtraFormatter`) - Claude proposed the custom log formatter that appends `extra={}` fields to log lines. The team tested and deployed it.

All other modules (`ai_service.py`, `cache_store.py`, `researcher.py`, `cli.py`, `app.py`) were written by the team with standard IDE autocompletion assistance.

Libraries: `tenacity` (retries), `pydantic-settings` (config), `httpx` (async HTTP), `pytest-asyncio` (async tests), `streamlit` (web UI).

References

- [1] Anthropic API documentation. <https://docs.anthropic.com>
- [2] `tenacity` retry library. <https://tenacity.readthedocs.io>

- [3] pydantic-settings documentation. https://docs.pydantic.dev/latest/concepts/pydantic_settings/
- [4] Python asyncio documentation: asyncio.timeout, asyncio.gather. <https://docs.python.org/3/library/asyncio.html>
- [5] Tavily AI Search API. <https://docs.tavily.com>
- [6] Google Gemini API. <https://ai.google.dev/gemini-api/docs>

A. Appendix - Reproducing the benchmark

```
git clone https://github.com/shahin1717/researchassistant
cd researchassistant
python -m venv .venv && source .venv/bin/activate
pip install -r requirements.txt
cp .env.example .env # add GOOGLE_API_KEY and TAVILY_API_KEY
python scripts/bench.py
```

The benchmark runs 5 questions sequentially (one source at a time) and then in parallel (`asyncio.gather`), printing wall-clock seconds and speedup per question. Numbers in the report were measured on a MacBook Pro M1 with a stable broadband connection. Provider latency varies; expect $\pm 20\%$ variance.

End of report - thank you for reading.